

# TSA-2018. Module 2. Time Series Regression

*N. Balakrishna and T. V. Ramanathan*

*August 2018*

## Modeling Time Series with Trend

The observed time series may be generated by a regression type model, in which certain functions of time are taken as independent variables. For example, a series with trend is shown below:

$$X_t = m_t + E_t. \quad (1)$$

The function  $m_t$  may be linear or non-linear in time  $t$ .

For example,  $m_t$  may be a polynomial in  $t$  of the type:

$$m_t = \beta_0 + \beta_1 t + \beta_2 \frac{t^2}{2!} + \beta_3 \frac{t^3}{3!} + \dots + \beta_p \frac{t^p}{p!} \quad (2)$$

## Method of Least Squares

In this method we fit a parametric family of functions for  $m_t$  by minimizing the sum of squares due to error, i.e.,

$$\sum_{t=1}^n (X_t - m_t)^2.$$

The order  $p$  of the polynomial is a *model selection* problem. That is, we fit polynomials of degrees  $1, 2, \dots$  for  $m_t$  and choose the one for which the least squares method provides a close fit for the data.

Let the estimates of the coefficients be denoted by  $\hat{\beta}_i, i = 0, 1, 2, \dots, p$ .

The estimate of the trend function is then

$$\hat{m}_t = \beta_0 + \hat{\beta}_1 t + \hat{\beta}_2 \frac{t^2}{2!} + \hat{\beta}_3 \frac{t^3}{3!} + \dots + \hat{\beta}_p \frac{t^p}{p!}.$$

Now compute the residuals:

$$\hat{E}_t = X_t - \hat{m}_t, t = 1, 2, \dots, n$$

and test for stationarity of  $\{\hat{E}_t\}$ .

## Global Temperature Deviations

Global mean land-ocean temperature deviations (from 1951-1980 average), measured in degrees centigrade.

```
library(astsa)
data(globtemp)
#help(globtemp)
globtemp=ts(globtemp) #save as a ts object
length(globtemp)
```

```
## [1] 136
```

### Outline of the Approach:

- Select the first  $n_{fit}$  observations as the Fitting/Training portion.
- Select the next  $L=n_{fore}$  observations as the Holdout/Test portion.
- Fit MLR models to the Fitting portion of the data.
- Assess the fits. Compute in-sample model selection criteria.
- Use the fitted models to forecast/predict the response time series for the Holdout portion, using coefficients from the fitted models and values of independent variables from the hold-out portion.
- Use forecast evaluation criteria as out-of-sample model selection criteria.
- Use the in-sample and out-of-sample criteria to select the best model(s).
- Use the best model(s) to forecast the future responses.

```
# Fitting Portion: globtemp from 1880:1993
xfit=globtemp[11:104]
nfit=length(xfit) #94
```

```
# Holdout Portion: globtemp from 1994:2003
xfore=gtemp[105:114]
nfore=length(xfore) #10
```

```
# k=1: Linear Trend Model fit to globtemp from 1880:2003, the
# fitting portion of the data.
# Set up a time vector as an independent variable
tfit=time(xfit) # time is from 1:94
```

```
# OLS fit: Linear Trend Model
mlr.lin = lm(xfit~tfit)
summary(mlr.lin)
```

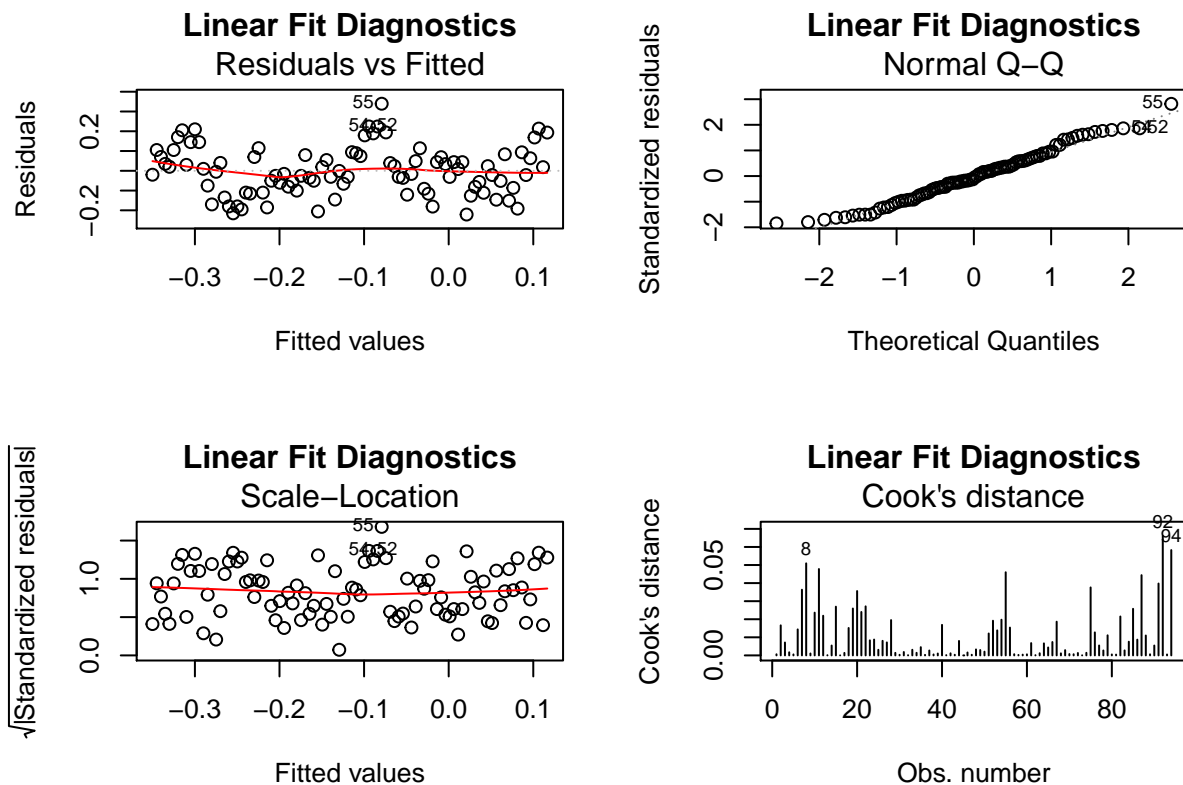
```
##
## Call:
## lm(formula = xfit ~ tfit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.22120 -0.08511 -0.01026  0.07323  0.33917
##
## Coefficients:
```

```
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.3551819  0.0251696  -14.11  <2e-16 ***
## tfit        0.0050184  0.0004601   10.91  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.121 on 92 degrees of freedom
## Multiple R-squared:  0.5639, Adjusted R-squared:  0.5592
## F-statistic: 119 on 1 and 92 DF,  p-value: < 2.2e-16
```

`anova(mlr.lin)`

```
## Analysis of Variance Table
##
## Response: xfit
##           Df Sum Sq Mean Sq F value    Pr(>F)
## tfit        1  1.7429   1.74294    118.96 < 2.2e-16 ***
## Residuals   92  1.3479   0.01465
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Residual Diagnostic Plots for mlr.lin
windows()
par(mfrow=c(2,2)) # plot 4 figures, 2 in each of 2 rows
plot(mlr.lin, main="Linear Fit Diagnostics", which = 1:4)
```



```

# k=2: Quadratic Trend Model fit
# Set up quadratic term independent variable
tsqfit=tfit^2/factorial(2)

mlr.quad=lm(xfit~tfit+tsqfit)
summary(mlr.quad)

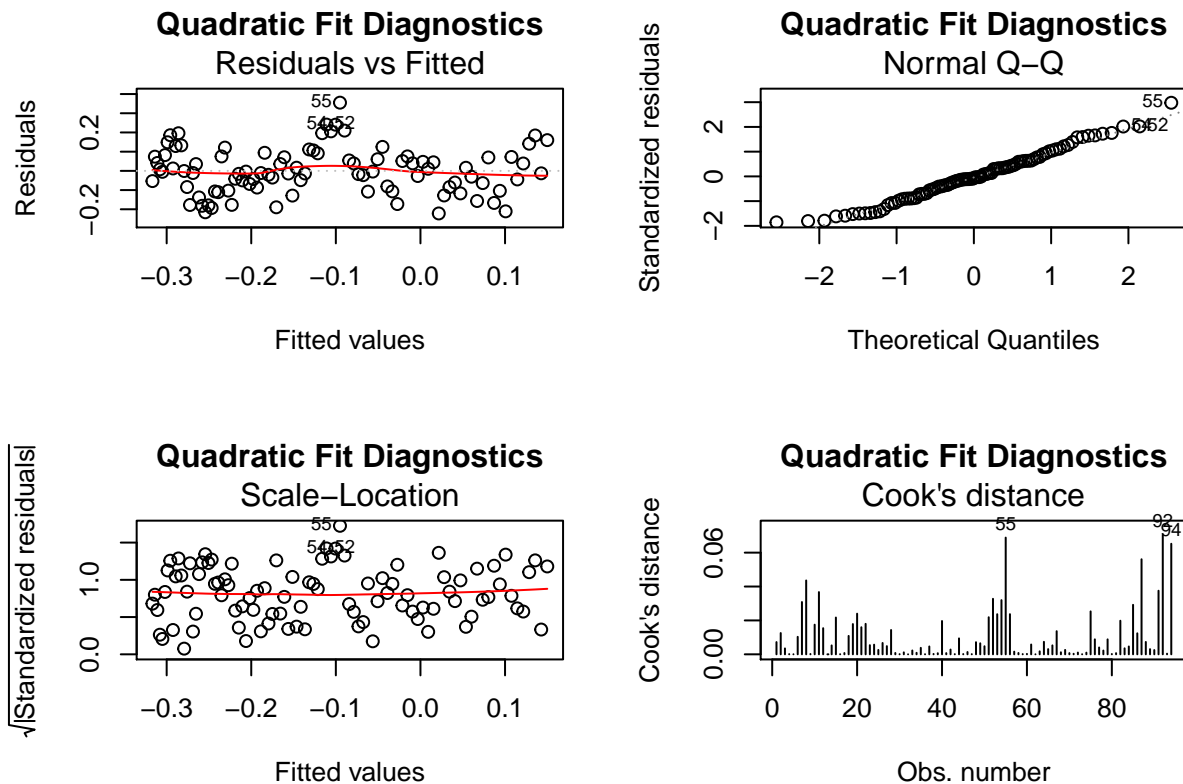
##
## Call:
## lm(formula = xfit ~ tfit + tsqfit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.22167 -0.08431 -0.00803  0.07353  0.35509
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.196e-01  3.815e-02  -8.376 6.33e-13 ***
## tfit         2.794e-03  1.854e-03   1.507   0.135
## tsqfit       4.683e-05  3.781e-05   1.238   0.219
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1207 on 91 degrees of freedom
## Multiple R-squared:  0.5711, Adjusted R-squared:  0.5617
## F-statistic: 60.59 on 2 and 91 DF,  p-value: < 2.2e-16

anova(mlr.quad)

## Analysis of Variance Table
##
## Response: xfit
##           Df Sum Sq Mean Sq F value Pr(>F)
## tfit       1 1.74294  1.74294 119.6521 <2e-16 ***
## tsqfit     1 0.02234  0.02234   1.5334 0.2188
## Residuals 91 1.32557  0.01457
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# Residual Diagnostic Plots for mlr.quad
windows()
par(mfrow=c(2,2))
plot(mlr.quad, main="Quadratic Fit Diagnostics",which = 1:4)

```



```
# k=3: Cubic Trend Model fit
# Set up cubic term independent variable
tcubfit=tfrit~3/factorial(3)

mlr.cub=lm(xfit~tfrit+tsqfit+tcubfit)
summary(mlr.cub)

##
## Call:
## lm(formula = xfit ~ tfrit + tsqfit + tcubfit)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.23823 -0.08700 -0.00369  0.06860  0.34501
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2.727e-01  5.162e-02  -5.283 8.77e-07 ***
## tfrit        -2.972e-03  4.681e-03  -0.635  0.527
## tsqfit        3.487e-04  2.283e-04   1.527  0.130
## tcubfit      -6.355e-06  4.741e-06  -1.340  0.183
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1202 on 90 degrees of freedom
## Multiple R-squared:  0.5795, Adjusted R-squared:  0.5655
```

```
## F-statistic: 41.35 on 3 and 90 DF, p-value: < 2.2e-16
```

```
anova(mlr.cub)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: xfit
```

```
##          Df Sum Sq Mean Sq F value Pr(>F)
## tfit      1 1.74294 1.74294 120.6995 <2e-16 ***
## tsqfit    1 0.02234 0.02234   1.5469 0.2168
## tcubfit    1 0.02594 0.02594   1.7966 0.1835
## Residuals 90 1.29963 0.01444
```

```
## ---
```

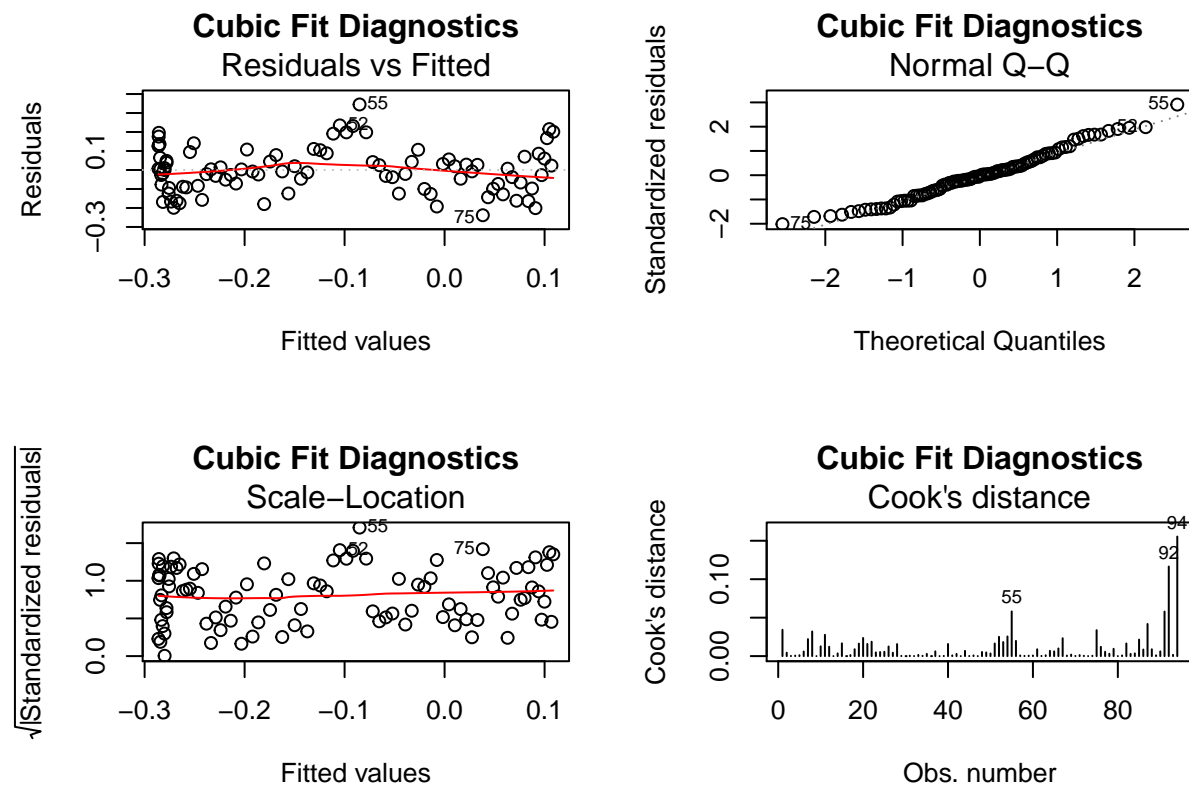
```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# Residual Diagnostic Plots for mlr.cub
```

```
windows()
```

```
par(mfrow=c(2,2))
```

```
plot(mlr.cub, main="Cubic Fit Diagnostics",which = 1:4)
```



```
# Plot time series with Fits from k=1,2,3 Models
```

```
par(mfrow=c(2,2))
```

```
ts.plot(xfit) # Time Series Plot
```

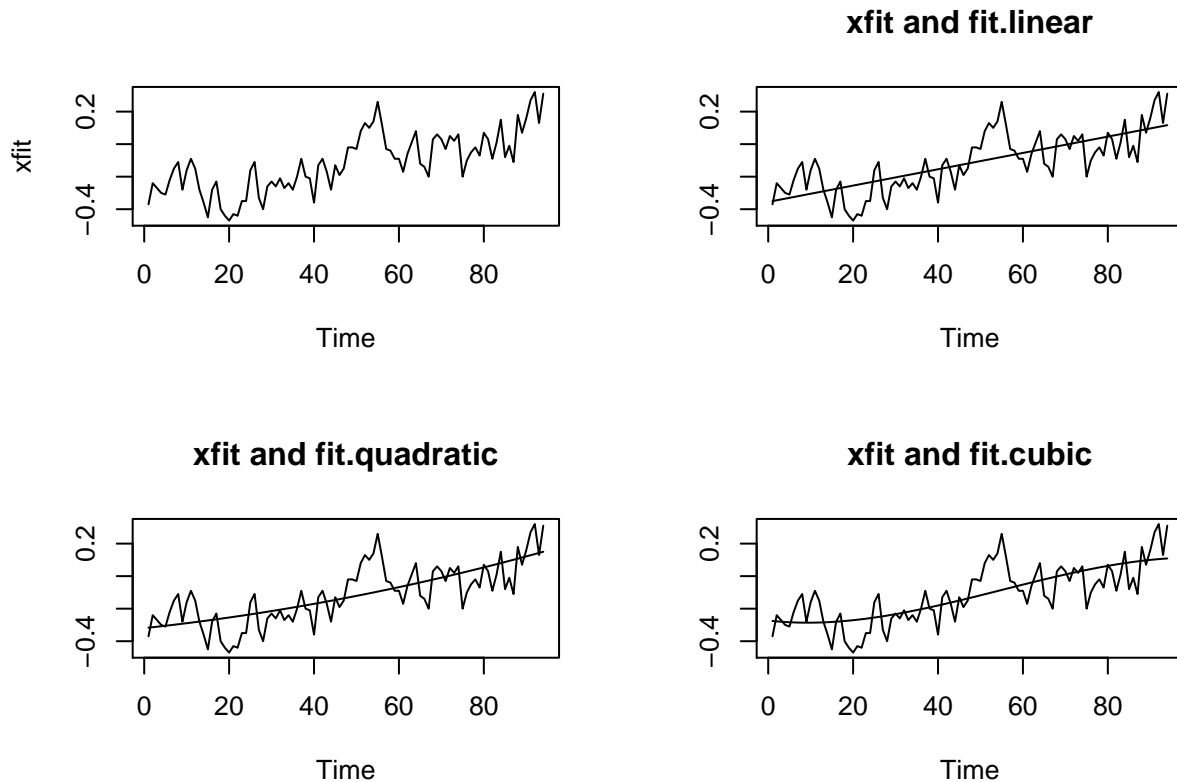
```
# Plot of xfit vs mlr.lin$fitted
```

```
plin=cbind(xfit,mlr.lin$fitted)
```

```
ts.plot(plin,main="xfit and fit.linear")
```

```
pquad=cbind(xfit,mlr.quad$fitted)
```

```
ts.plot(pquad,main="xfit and fit.quadratic")
pcub=cbind(xfit,mlr.cub$fitted)
ts.plot(pcub,main="xfit and fit.cubic")
```



### In-sample Model selection criteria:

Which model is better? Compare residual mean squares in ANOVA tables, AIC, BIC. The smaller, the better.

```
#Residual Mean Squares
sigsq.lin=anova(mlr.lin)[["Mean Sq"]][2]
sigsq.quad=anova(mlr.quad)[["Mean Sq"]][3]
sigsq.cub=anova(mlr.cub)[["Mean Sq"]][4]
```

```
### Linear Trend Model, k=1
k=1
# Akaike Information Criterion, AIC
AIC.lin = AIC(mlr.lin)/nfit
cat("AIC for k=1 is",AIC.lin, "\n")
```

```
## AIC for k=1 is -1.343035
```

```
# Or RHS is AIC(mlr.lin)/length(xfit)
# Or RHS is AIC(logLik(mlr.lin))/nfit
# Or RHS is log(sigsq.lin)+(nfit+2*k)/nfit
```

```

# Bayesian Information Criterion, BIC
BIC.lin = BIC(mlr.lin)/nfit
cat("BIC for k=1 is", BIC.lin, "\n")

## BIC for k=1 is -1.261866

# Or RHS is BIC(mlr.lin)/length(xfit)
# Or RHS is BIC(logLik(mlr.lin))/nfit
# Or RHS is log(sigsq.lin)+(k*log(nfit))/nfit

# Corrected AIC, i.e., AICc using formula
AICc.lin=log(sigsq.lin)+(nfit+k)/(nfit-k-2)
cat("AICc for k=1 is", AICc.lin, "\n")

## AICc for k=1 is -3.17928

### Quadratic Trend Model, k=2
k=2
# AIC: Try the other ways
AIC.quad = AIC(mlr.quad)/nfit
cat("AIC for k=2 is", AIC.quad, "\n")

## AIC for k=2 is -1.338469

# BIC: Try the other ways
BIC.quad = BIC(mlr.quad)/nfit
cat("BIC for k=2 is", BIC.quad, "\n")

## BIC for k=2 is -1.230244

# AICc using formula
AICc.quad=log(sigsq.quad)+(nfit+k)/(nfit-k-2)
cat("AICc for k=2 is", AICc.quad, "\n")

## AICc for k=2 is -3.162351

### Cubic Trend Model, k=3
k=3
# AIC: Try the other ways
AIC.cub = AIC(mlr.cub)/nfit
cat("AIC for k=3 is", AIC.cub, "\n")

## AIC for k=3 is -1.336958

# BIC: Try the other ways
BIC.cub = BIC(mlr.cub)/nfit
cat("BIC for k=3 is", BIC.cub, "\n")

## BIC for k=3 is -1.201677

# AICc using formula
AICc.cub=log(sigsq.cub)+(nfit+k)/(nfit-k-2)
cat("AICc for k=3 is", AICc.cub, "\n")

## AICc for k=3 is -3.147845

```



## Forecasting Holdout/Test Data

```
# (iv) Forecast Evaluation for Different Models
# Set up and predict for next L=nfore=10 times in the future.
# That is, predict the Test/Forecast Validation portion
# of data from 1994 to 2003

# Set up for SLR (k=1) model
new <- data.frame(tfit=c(95:104))
# use predict function
pfore.lin=predict(mlr.lin,new,se.fit = TRUE)
pfore.lin$fit      #point predictions

##          1          2          3          4          5          6          7
## 0.1215649 0.1265832 0.1316016 0.1366200 0.1416384 0.1466568 0.1516752
##          8          9         10
## 0.1566936 0.1617120 0.1667303

#Note: you can also get all this info, if you like
pfore.lin$se.fit

##          1          2          3          4          5          6
## 0.02516961 0.02557015 0.02597266 0.02637705 0.02678324 0.02719115
##          7          8          9         10
## 0.02760070 0.02801183 0.02842445 0.02883851

pfore.lin$df

## [1] 92

pfore.lin$residual.scale

## [1] 0.121042

# Observed data for 1994-2003 is in xfore
# Forecast/Prediction errors: Observed - Predicted
efore.lin=xfore-pfore.lin$fit      #length is nfore=length(xfore)
```

## Forecast Evaluation Criteria:

Suppose we denote the forecast errors by  $e_n(\ell)$  where  $n$  is the forecast origin and  $\ell$  is the forecast lead.

We compute and discuss Mean error, Mean Percent Error, Mean Squared Error, Mean Absolute Error, Mean Absolute Percent Error.

$$ME = \bar{e} = \frac{1}{L} \sum_{\ell=1}^L e_n(\ell); \quad (3)$$

$$MPE = \frac{100}{L} \sum_{\ell=1}^L \{e_n(\ell)/X_{n+\ell}\}; \quad (4)$$

$$MSE = \frac{1}{L} \sum_{\ell=1}^L e_n^2(\ell); \quad (5)$$

$$MAE = \frac{1}{L} \sum_{\ell=1}^L |e_n(\ell)|; \quad (6)$$

$$MAPE = \frac{100}{L} \sum_{\ell=1}^L |e_n(\ell)/X_{n+\ell}|; \quad (7)$$

The first two criteria are Mean Error (ME) and Mean Percent Error (MPE) and indicate bias in forecasts.

The last three are Mean Squared Error (MSE), Mean Absolute Error (MAE) and Mean Absolute Percent Error (MAPE) and assess forecast accuracy.

In many cases, the MAPE is used for model selection, because it is a percentage which accounts for the scale of the data and can be useful for broad comparison across models/data. However, when data exhibits periodic behavior or some of the data values are very close to zero, the MAPE can be misleading. It is a good idea to explore the forecasts using plots and perhaps use other forecast accuracy criteria such as MSE or MAE.

In some situations, the information criteria and forecast accuracy criteria may select different models! The user relies on whichever is more important, in-sample accuracy or out-of-sample forecast accuracy.

```
# Forecast evaluation criteria for k=1 (linear)
me.lin=mean(efore.lin)           # Mean Error
mpe.lin=100*(mean(efore.lin/xfore)) # Mean Percent Error
mse.lin=sum(efore.lin**2)/nfore   # Mean Squared Error
mae.lin=mean(abs(efore.lin))      # Mean Absolute Error
mape.lin=100*mean(abs(efore.lin/xfore)) # Mean Absolute Percent Error
cat("Mean Error (ME) for k=1 is",me.lin, "\n")
```

```
## Mean Error (ME) for k=1 is 0.0548524
```

```
cat("Mean Percent Error (MPE) for k=1 is",mpe.lin, "\n")
```

```
## Mean Percent Error (MPE) for k=1 is -2.210729
```

```
cat("Mean Squared Error (MSE) for k=1 is",mse.lin, "\n")
```

```
## Mean Squared Error (MSE) for k=1 is 0.01405649
```

```
cat("Mean Absolute Error (MAE) for k=1 is",mae.lin, "\n")
```

```
## Mean Absolute Error (MAE) for k=1 is 0.09449081
```

```
cat("Mean Absolute Percent Error (MAPE) for k=1 is",mape.lin, "\n")
```

```
## Mean Absolute Percent Error (MAPE) for k=1 is 49.9734
```

```

# Set up for prediction with k=2 model
tfit=c(95:104)
tsqfit=tfit^2/factorial(2)
matq=matrix(c(tfit,tsqfit),nrow=10,ncol=2,dimnames=list(c(),c("tfit","tsqfit")))
matq

##          tfit tsqfit
## [1,]    95 4512.5
## [2,]    96 4608.0
## [3,]    97 4704.5
## [4,]    98 4802.0
## [5,]    99 4900.5
## [6,]   100 5000.0
## [7,]   101 5100.5
## [8,]   102 5202.0
## [9,]   103 5304.5
## [10,]  104 5408.0

newnq <- data.frame(matq)

pfore.quad=predict(mlr.quad,newnq,se.fit=TRUE)
pfore.quad$fit # point predictions

##          1          2          3          4          5          6          7
## 0.1571524 0.1644185 0.1717313 0.1790910 0.1864975 0.1939508 0.2014509
##          8          9         10
## 0.2089979 0.2165917 0.2242324

efore.quad=xfore-pfore.quad$fit # Forecast errors
me.quad=mean(efore.quad)
mpe.quad=100*(mean(efore.quad/xfore))
mse.quad=sum(efore.quad**2)/nfore
mae.quad=mean(abs(efore.quad))
mape.quad=100*mean(abs(efore.quad/xfore))

cat("Mean Error (ME) for k=2 is",me.quad, "\n")

## Mean Error (ME) for k=2 is 0.008588562
cat("Mean Percent Error (MPE) for k=2 is",mpe.quad, "\n")

## Mean Percent Error (MPE) for k=2 is -34.48205
cat("Mean Squared Error (MSE) for k=2 is",mse.quad, "\n")

## Mean Squared Error (MSE) for k=2 is 0.0108592
cat("Mean Absolute Error (MAE) for k=2 is",mae.quad, "\n")

## Mean Absolute Error (MAE) for k=2 is 0.09420397
cat("Mean Absolute Percent Error (MAPE) for k=2 is",mape.quad, "\n")

## Mean Absolute Percent Error (MAPE) for k=2 is 65.84175

# Set up for prediction with k=3 model
tfit=c(95:104)
tsqfit=tfit^2/factorial(2)

```

```
tcubfit=tfit^3/factorial(3)
matc=matrix(c(tfit,tsqfit,tcubfit),nrow=10,ncol=3,dimnames = list(c(),c("tfit","tsqfit","tcubfit")))
matc
```

```
##      tfit tsqfit tcubfit
## [1,]   95 4512.5 142895.8
## [2,]   96 4608.0 147456.0
## [3,]   97 4704.5 152112.2
## [4,]   98 4802.0 156865.3
## [5,]   99 4900.5 161716.5
## [6,]  100 5000.0 166666.7
## [7,]  101 5100.5 171716.8
## [8,]  102 5202.0 176868.0
## [9,]  103 5304.5 182121.2
## [10,] 104 5408.0 187477.3
```

```
newnc=data.frame(matc)
```

```
pfore.cub=predict(mlr.cub,newnc,se.fit = TRUE)
pfore.cub$fit # point predictions
```

```
##      1      2      3      4      5      6      7
## 0.1103007 0.1116486 0.1127351 0.1135539 0.1140985 0.1143626 0.1143400
##      8      9     10
## 0.1140241 0.1134087 0.1124874
```

```
efore.cub=xfore-pfore.cub$fit # Forecast errors
me.cub=mean(efore.cub)
mpe.cub=100*(mean(efore.cub/xfore))
mse.cub=sum(efore.cub**2)/nfore
mae.cub=mean(abs(efore.cub))
mape.cub=100*mean(abs(efore.cub/xfore))
```

```
cat("Mean Error (ME) for k=3 is",me.cub, "\n")
```

```
## Mean Error (ME) for k=3 is 0.08590403
```

```
cat("Mean Percent Error (MPE) for k=3 is",mpe.cub, "\n")
```

```
## Mean Percent Error (MPE) for k=3 is 17.5012
```

```
cat("Mean Squared Error (MSE) for k=3 is",mse.cub, "\n")
```

```
## Mean Squared Error (MSE) for k=3 is 0.01912071
```

```
cat("Mean Absolute Error (MAE) for k=3 is",mae.cub, "\n")
```

```
## Mean Absolute Error (MAE) for k=3 is 0.1022939
```

```
cat("Mean Absolute Percent Error (MAPE) for k=3 is",mape.cub, "\n")
```

```
## Mean Absolute Percent Error (MAPE) for k=3 is 46.67193
```

## Comparison of forecasted and actual values on Holdout Data

```
### Plot of time series with fits and out of sample forecasts
### Linear Trend Model fits and forecasts
linff=c(mlr.lin$fitted,pfore.lin$fit)
```



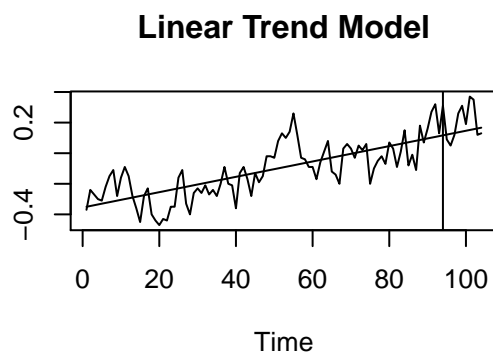
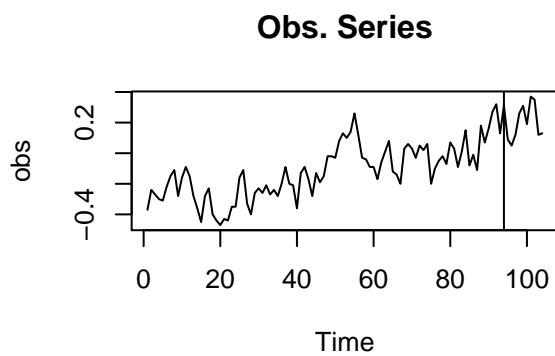
```

### Quadratic Trend Model fits and forecasts
quadff=c(mlr.quad$fitted,pfore.quad$fit)
### Cubic Trend Model fits and forecasts
cubff=c(mlr.cub$fitted,pfore.cub$fit)

# observed data (n=98)
obs=c(xfit,xfore)
# Bind observed data and fits+forecasts
obslin=cbind(obs,linff)
obsquad=cbind(obs,quadff)
obsclub=cbind(obs,cubff)
time=c(1:length(obs))

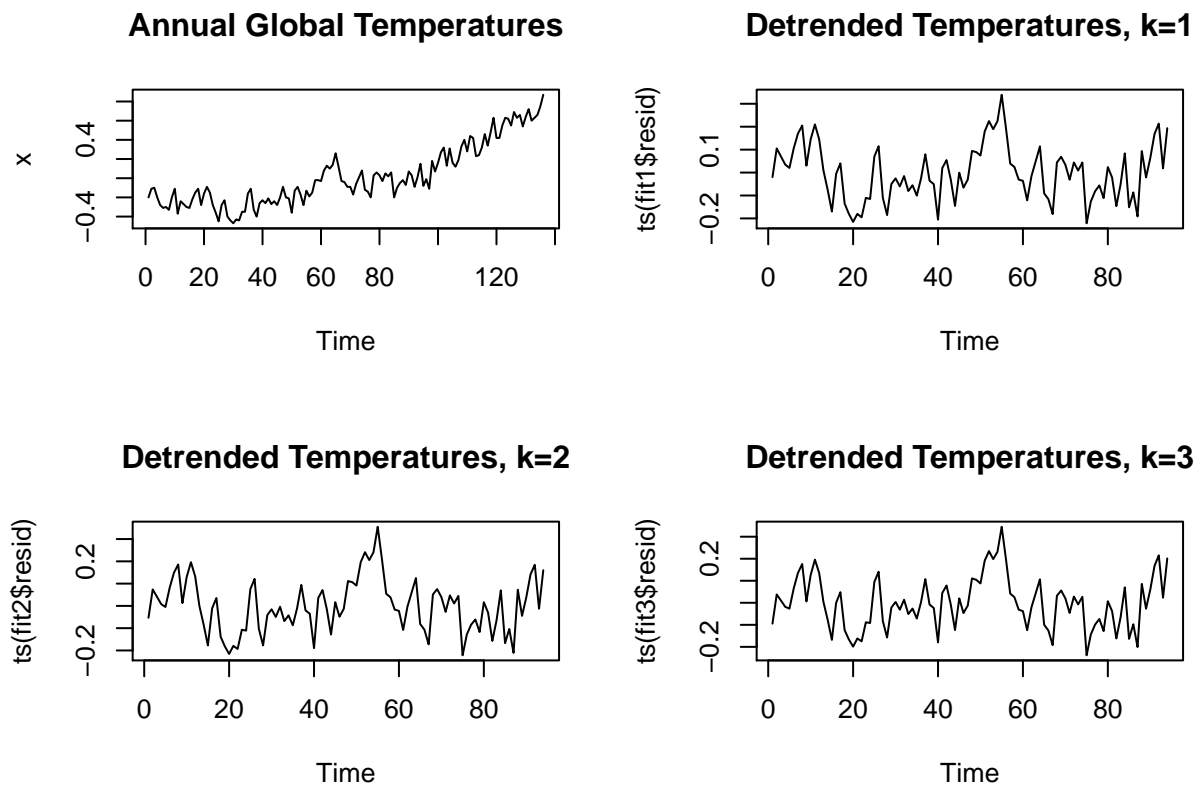
par(mfrow=c(2,2))
ts.plot(obs,main="Obs. Series")
abline(v=94)
ts.plot(obslin,main="Linear Trend Model")
abline(v=94)
ts.plot(obsquad,main="Quadratic Trend Model")
abline(v=94)
ts.plot(obsclub,main="Cubic Trend Model")
abline(v=94)

```



## Detrending the Time Series

```
### Detrending a Time Series after a Polynomial Model Fit
x=globtemp
fit1=mlr.lin
fit2=mlr.quad
fit3=mlr.cub
par(mfrow=c(2,2))
### Original series as x
ts.plot(x,main="Annual Global Temperatures")
### Detrended series, k=1
ts.plot(ts(fit1$resid),main="Detrended Temperatures, k=1")
### Detrended series, k=2
ts.plot(ts(fit2$resid),main="Detrended Temperatures, k=2")
### Detrended series, k=3
ts.plot(ts(fit3$resid),main="Detrended Temperatures, k=3")
```



## Modeling Time Series with Trend and Seasonality

A time series  $X_t$  may have an additive decomposition:

$$X_t = m_t + S_t + E_t, \quad (8)$$

where the trend component could be a polynomial of degree  $p$  (as before) and  $S_t$  may be a periodic function with period  $s$ .

$S_t$  may be modeled using **indicator functions of time**:

$$S_t = \sum_{i=1}^s \gamma_i IND_{i,t}$$

where  $IND_{i,t}$  are indicator functions of time  $t$ , defined by

$$IND_{i,t} = \begin{cases} 1 & \text{if } t \text{ corresponds to seasonal period } i \\ 0 & \text{otherwise} \end{cases}$$

When  $s = 12$ , we set up indicator variables, one for each month. That is, for  $i = 1, \dots, s$ , define

For any Jan. we have  $IND_{1,t} = 1$ , while  $IND_{i,t} = 0$  for  $i \neq 1$

For any Feb. we have  $IND_{2,t} = 1$ , while  $IND_{i,t} = 0$  for  $i \neq 2$

etc.


For any Dec. we have  $IND_{12,t} = 1$ , while  $IND_{i,t} = 0$  for  $i \neq 12$

### Example: Accidental Deaths Data

```
accd <- read.csv("C:/Users/STATISTICS/Desktop/Manipal/Data/accdeaths.csv",header=F)

accd=ts(accd)    #length =72
ts.plot(accd,main="Accidental Deaths")
```

## Accidental Deaths



```
# Regression with Polynomial terms for m_t and
# Monthly indicators for S_t
# Polynomial terms: t, t^2/2!
t=1:length(accd)
tsq=t^2/factorial(2)
# Seasonal terms: 12 monthly indicators
per=12
sets=length(accd)/per
month=factor(rep(1:per,sets))

### Model 1. t,tsq,12 seasonal indicators,no intercept
accd.model1 = lm(accd~t+tsq+month -1)
summary(accd.model1)
```

```
##
## Call:
## lm(formula = accd ~ t + tsq + month - 1)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-587.84	-153.18	-9.46	178.04	609.04

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
t	-71.9782	6.3391	-11.36	2.27e-16 ***
tsq	1.6531	0.1682	9.83	5.83e-14 ***



```

## month1    9133.8707    143.3724    63.71 < 2e-16 ***
## month2    8393.6101    144.3576    58.15 < 2e-16 ***
## month3    9191.8631    145.2706    63.27 < 2e-16 ***
## month4    9409.4631    146.1104    64.40 < 2e-16 ***
## month5   10285.7432    146.8766    70.03 < 2e-16 ***
## month6   10768.2037    147.5690    72.97 < 2e-16 ***
## month7   11637.3443    148.1883    78.53 < 2e-16 ***
## month8   10943.6653    148.7352    73.58 < 2e-16 ***
## month9    9903.1664    149.2112    66.37 < 2e-16 ***
## month10  10194.1812    149.6181    68.14 < 2e-16 ***
## month11   9681.7095    149.9582    64.56 < 2e-16 ***
## month12   9938.5848    150.2340    66.15 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 275.4 on 58 degrees of freedom
## Multiple R-squared:  0.9992, Adjusted R-squared:  0.999
## F-statistic: 5292 on 14 and 58 DF, p-value: < 2.2e-16
### Model 2. t,tsq,11 seasonal indicators and intercept
mon=factor(rep(c(1:(per-1),0),sets))
accd.model2=lm(accd~t+tsq+mon)
summary(accd.model2)

##
## Call:
## lm(formula = accd ~ t + tsq + mon)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -587.84 -153.18   -9.46  178.04  609.04
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  9938.5848   150.2340   66.154 < 2e-16 ***
## t            -71.9782     6.3391  -11.355 2.27e-16 ***
## tsq             1.6531     0.1682    9.830 5.83e-14 ***
## mon1         -804.7141   159.9783   -5.030 5.04e-06 ***
## mon2        -1544.9747   159.8158   -9.667 1.07e-13 ***
## mon3         -746.7217   159.6715   -4.677 1.79e-05 ***
## mon4         -529.1217   159.5435   -3.316 0.00158 **
## mon5          347.1584   159.4301    2.177 0.03352 *
## mon6          829.6189   159.3304    5.207 2.64e-06 ***
## mon7         1698.7595   159.2438   10.668 2.68e-15 ***
## mon8         1005.0805   159.1703    6.314 4.11e-08 ***
## mon9          -35.4184   159.1105   -0.223 0.82463
## mon10         255.5964   159.0655    1.607 0.11352
## mon11        -256.8753   159.0368   -1.615 0.11170
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 275.4 on 58 degrees of freedom
## Multiple R-squared:  0.9325, Adjusted R-squared:  0.9174
## F-statistic: 61.65 on 13 and 58 DF, p-value: < 2.2e-16

```


```

#tq=t^3/factorial(3)
#accd.model3=lm(accd~t+tsq+tq+mon)
#AIC increased


### Plot data and fits from Model 1 and Model 2
### Verify Model 1 and Model 2 are equivalent!
windows()
par(mfrow=c(1,2))
ts.plot(accd,main="Accidental deaths, Model 1 Fit")
lines(accd.model1$fit,col="blue", lty="dashed")
ts.plot(accd,main="Accidental deaths, Model 2 Fit")
lines(accd.model2$fit,col="red", lty="dotted")

```

**Accidental deaths, Model 1 Fit**



**Accidental deaths, Model 2 Fit**



**Exercise.** Fit a suitable time series regression model to Air Passengers data and use this to forecast the next year.